

# Computer Networking-lab1-Repot

课程名称：计算机网络 任课教师：田臣/李文中

学院	Dept. of Computer Science and Technology	专业（方向）	CS
学号	181830044	姓名	董宸郅
Email	pdon#foxmail.com	开始/完成日期	3.19/3.25

## 实验名称：Leaning Switch

### 实验目的

- 深入理解`switch`的工作原理
- 理解并实现具有“学习能力”的switch
- 进一步熟悉switchyard的代码接口和整体框架
- 加强在实验环境下调试代码的能力

### 实验内容

#### 理论知识

##### Switch概念解析

Switch基于mac地址识别，能完成封装转发数据功能的设备。可以记录一定数量的mac地址，放在**内部地址表**中，通过在数据帧的始发者和接收者之间建立临时的交换路径，使数据从源地址到达目的地址。（与之相比，Hub是一种共享设备，本身不能识别目的地址。采用广播的形式传输数据，即向所有端口传送数据。）

##### ARP

地址解析协议(Address Resolution Protocol)，是根据IP地址获取物理地址的一个TCP/IP协议。主机发送信息时将**包含目标IP地址的ARP请求广播到局域网络上的所有主机**，并接收返回消息，以此确定目标的物理地址；收到返回消息后将该IP地址和物理地址**存入本机ARP缓存中并保留一定时间**，下次请求时直接查询ARP缓存以节约资源。

### 实验步骤（含结果与关键代码）

**注：**

- ①对于之后几个Task中与Task 2基本一致的代码不再专门说明展示，将只说明其特有的关键代码
- ②Deploy阶段的mininet拓扑结构即为默认提供的结构：中央的switch分别连接server1、server2、client。

#### Task 2: Basic Switch

Coding

```

1 # add an empty heapq to store forwarding rules
2 # each element of tab should be a tuple like: (traffic, host, intf) --
  >traffic means the number of packets related
3 # when a heapq consists of tuples, it is organized based on the tuples' first
  elements
4 tab = []

```

```

1 # recording section
2 # src already recorded
3 if packet[0].src in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].src:
8             cur = rule
9             break
10    # first delete, then add a new one
11    new_rule = (cur[0], cur[1], input_port)
12    tab.remove(cur)
13    tab.append(new_rule)
14    heapq.heapify(tab)
15 # src not recorded yet
16 else:
17     # forwarding table is full
18     if len(tab) == max_rules:
19         heapq.heappop(tab)
20     tab.append((0, packet[0].src, input_port))
21     heapq.heapify(tab)

```

```

1 # forwarding section
2 # dst already recorded
3 if packet[0].dst in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].dst:
8             cur = rule
9             break
10    log_debug ("Flooding packet {} to {}".format(packet, cur[2]))
11    net.send_packet(cur[2], packet)
12    new_rule = (cur[0]+1, cur[1], cur[2])
13    tab.remove(cur)
14    tab.append(new_rule)
15    heapq.heapify(tab)
16 # dst not recorded yet
17 else:
18     # do nothing if dst is the switch itself
19     if packet[0].dst not in mymacs:
20         for intf in my_interfaces:
21             if input_port != intf.name:
22                 log_debug ("Flooding packet {} to {}".format(packet,
23                     intf.name))
23                 net.send_packet(intf.name, packet)

```

## Deploying

## server1:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	30:00:00:00:00...	Broadcast	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
2	0.116833379	Private_00:00:00...	30:00:00:00:00...	ARP	42	192.168.100.1 is at 10:00:00:00:00:01
3	0.669877596	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0bec, seq=1/256, ttl=64 (reply in 4)
4	0.793030851	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0bec, seq=1/256, ttl=64 (request in 3)
5	0.996696301	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0bec, seq=2/512, ttl=64 (reply in 6)
6	1.096822978	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0bec, seq=2/512, ttl=64 (request in 5)
7	5.934673661	Private_00:00:00...	30:00:00:00:00...	ARP	42	Who has 192.168.100.3? Tell 192.168.100.1
8	6.396574253	30:00:00:00:00...	Private_00:00:00...	ARP	42	192.168.100.3 is at 30:00:00:00:00:01

## server2:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	30:00:00:00:00...	Broadcast	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3

在client的cli中输入以下命令，以向server1发出两次请求：

```
1 | ping -c 2 192.168.100.1
```

根据上述截图，由于server2有且仅有一次收到ARP包询问是否拥有server1的IP地址，可知第一次请求时switch不知道server1的mac地址，第二次则已完成记录，直接发给server1，验证了basic learning的逻辑。

## Task 3: Timeouts

### Coding

```
1 | # add an empty dict intended for host-(intf,last_time) pairs
2 | tab={}
```

```
1 | # when new packet comes, record its information or update recorded
   | information
2 | # since 'dictionary' is used and its keys(hosts' mac addresses) are all
   | unique, the following line can do it all
3 | tab[packet[0].src] = {'intf': input_port, 'last': time.time()}
4 | # delete timeout entries in the forwarding table
5 | for host in list(tab):
6 |     if time.time()-tab[host]['last'] > 10:
7 |         del tab[host]
```

### Testing

Results for test scenario switch tests: 9 passed, 0 failed, 0 pending

Passed:

- 1 An Ethernet frame with a broadcast destination address should arrive on eth1
- 2 The Ethernet frame with a broadcast destination address should be forwarded out ports eth0 and eth2
- 3 An Ethernet frame from 20:00:00:00:00:01 to 30:00:00:00:00:02 should arrive on eth0
- 4 Ethernet frame destined for 30:00:00:00:00:02 should arrive on eth1 after self-learning
- 5 Timeout for 20s
- 6 An Ethernet frame from 20:00:00:00:00:01 to 30:00:00:00:00:02 should arrive on eth0
- 7 Ethernet frame destined for 30:00:00:00:00:02 should be flooded out eth1 and eth2
- 8 An Ethernet frame should arrive on eth2 with destination address the same as eth2's MAC address
- 9 The hub should not do anything in response to a frame arriving with a destination address referring to the hub itself.

All tests passed!

## Deploying

server1:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.00000000	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d4f, seq=1/256, ttl=64 (reply in 2)
2	0.109098147	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0d4f, seq=1/256, ttl=64 (request in 1)
3	0.939345013	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d4f, seq=2/512, ttl=64 (reply in 4)
4	1.039499469	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0d4f, seq=2/512, ttl=64 (request in 3)
5	5.062906588	30:00:00:00:00:01	Private_00:00:01	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
6	5.134184208	Private_00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.1
7	5.163162672	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01
8	5.478011697	30:00:00:00:00:01	Private_00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01
9	53.643615899	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d53, seq=1/256, ttl=64 (reply in 10)
10	53.745554743	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0d53, seq=1/256, ttl=64 (request in 9)
11	54.604462003	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d53, seq=2/512, ttl=64 (reply in 12)
12	54.705857091	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0d53, seq=2/512, ttl=64 (request in 11)
13	58.804501298	30:00:00:00:00:01	Private_00:00:01	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
14	58.892697087	Private_00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.1
15	58.904925732	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01
16	59.222410076	30:00:00:00:00:01	Private_00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01

server2:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.00000000	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d4f, seq=1/256, ttl=64 (no response found!)
2	53.643611694	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0d53, seq=1/256, ttl=64 (no response found!)

在client的cli中输入以下命令，以向server1发出四次请求（分两次发送，且间隔时间大于10s）：

```
1 ping -c 2 192.168.100.1
2 # more than 10s later
3 ping -c 2 192.168.100.1
```

从上述截图看出，server2收到两次client对server1的请求。

第一次请求时switch不知道server1的地址，进行记录后第二次请求不需要再次广播。

第三次请求时，由于已经过了规定的10s，由timeout机制，记录被清除，所以再次广播，到了第四次，就和第二次一样，直接发给server1。

从而，验证了timeout逻辑。

## Task 4: Least Recently Used

Coding

```

1 # add an empty heapq to store forwarding rules
2 # each element of tab should be a tuple like: (last_time, host, intf)
3 # when a heapq consists of tuples, it is organized based on the tuples' first
  elements
4 # therefore the least recently used rule(its 'last_time' being the 'least'
  when treated as a number) is always at the front, which is what we expect
5 tab = []

```

```

1 # recording section
2 # src already recorded
3 if packet[0].src in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].src:
8             cur = rule
9             break
10    # adjust the heapq
11    # first delete, then add a new one since it's a tuple
12    tab.remove(cur)
13 # src not recorded yet
14 else:
15     # forwarding table is full
16     if len(tab) == max_rules:
17         heapq.heappop(tab)
18     tab.append((time.time(), packet[0].src, input_port))
19     heapq.heapify(tab)

```

```

1 # forwarding section
2 # dst already recorded
3 if packet[0].dst in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].dst:
8             cur = rule
9             break
10    log_debug ("Flooding packet {} to {}".format(packet, cur[2]))
11    net.send_packet(cur[2], packet)
12    # adjust the heapq
13    # first delete, then add a new one since it's a tuple
14    new_rule = (time.time(), cur[1], cur[2])
15    tab.remove(cur)
16    tab.append(new_rule)
17    heapq.heapify(tab)
18 # dst not recorded yet
19 else:
20     # do nothing if dst is the switch itself
21     if packet[0].dst not in mymacs:
22         for intf in my_interfaces:
23             if input_port != intf.name:
24                 log_debug ("Flooding packet {} to {}".format(packet,
25                     intf.name))
26                 net.send_packet(intf.name, packet)

```

## Testing

Passed:

- 1 An Ethernet frame with a broadcast destination address should arrive on eth1
- 2 The Ethernet frame with a broadcast destination address should be forwarded out ports eth0, eth2, eth3 and eth4
- 3 An Ethernet frame from 20:00:00:00:00:01 to 30:00:00:00:00:02 should arrive on eth0
- 4 Ethernet frame destined for 30:00:00:00:00:02 should arrive on eth1 after self-learning
- 5 An Ethernet frame from 20:00:00:00:00:03 to 30:00:00:00:00:02 should arrive on eth2
- 6 Ethernet frame destined for 30:00:00:00:00:02 should arrive on eth1 after self-learning
- 7 An Ethernet frame from 30:00:00:00:00:04 to 20:00:00:00:00:01 should arrive on eth3
- 8 Ethernet frame destined to 20:00:00:00:00:01 should arrive on eth0 after self-learning
- 9 An Ethernet frame from 20:00:00:00:00:01 to 30:00:00:00:00:04 should arrive on eth0
- 10 Ethernet frame destined to 20:00:00:00:00:01 should arrive on eth3 after self-learning
- 11 An Ethernet frame from 40:00:00:00:00:05 to 20:00:00:00:00:01 should arrive on eth4
- 12 Ethernet frame destined to 20:00:00:00:00:01 should arrive on eth0 after self-learning
- 13 An Ethernet frame from 30:00:00:00:00:05 to 20:00:00:00:00:01 should arrive on eth4
- 14 Ethernet frame destined to 20:00:00:00:00:01 should arrive on eth0 after self-learning
- 15 An Ethernet frame from 20:00:00:00:00:05 to 30:00:00:00:00:02 should arrive on eth4
- 16 Ethernet frame destined to 30:00:00:00:00:02 should be flooded to eth0, eth1, eth2 and eth3
- 17 An Ethernet frame should arrive on eth2 with destination address the same as eth2's MAC address
- 18 The hub should not do anything in response to a frame arriving with a destination address referring to the hub itself.

All tests passed!

## Deploying

server1:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e1d, seq=1/256, ttl=64 (reply in 2)
2	0.103000481	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e1d, seq=1/256, ttl=64 (request in 1)
3	0.941412336	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e1d, seq=2/512, ttl=64 (reply in 4)
4	1.041711176	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e1d, seq=2/512, ttl=64 (request in 3)
5	5.082535485	30:00:00:00:00:01	Private_00:00:01	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
6	5.182721656	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01
7	5.192113306	Private_00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.1
8	5.602972490	30:00:00:00:00:01	Private_00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01
9	9.099949357	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.2? Tell 192.168.100.3
10	29.928844407	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e20, seq=1/256, ttl=64 (reply in 11)
11	30.028946144	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e20, seq=1/256, ttl=64 (request in 10)
12	30.881218559	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e20, seq=2/512, ttl=64 (reply in 13)
13	30.982214852	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e20, seq=2/512, ttl=64 (request in 12)
14	35.113281397	30:00:00:00:00:01	Private_00:00:01	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
15	35.215979139	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01

## server2:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e1d, seq=1/256, ttl=64 (no response found!)
2	9.099949461	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.2? Tell 192.168.100.3
3	9.205943341	20:00:00:00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.2 is at 20:00:00:00:00:01
4	9.632413597	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x0e1e, seq=1/256, ttl=64 (reply in 5)
5	9.734988380	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e1e, seq=1/256, ttl=64 (request in 4)
6	10.070383175	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x0e1e, seq=2/512, ttl=64 (reply in 7)
7	10.170648244	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x0e1e, seq=2/512, ttl=64 (request in 6)
8	14.916069766	20:00:00:00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.2
9	15.316532769	30:00:00:00:00:01	20:00:00:00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01
10	29.928844717	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x0e20, seq=1/256, ttl=64 (no response found!)

在client的cli中输入以下命令：

```
1 ping -c 2 192.168.100.1
2 ping -c 2 192.168.100.2
3 ping -c 2 192.168.100.2
4 ping -c 2 192.168.100.1
```

首先，设置最多容纳的rule数目为2。

在最开始的client对server1的两次请求中，结果与之前task一致。

之后，进行若干次client对server2的请求，使server2的以达到清除有关server1对应端口信息的目的。

最后，再次进行client对server1的请求，server2的wireshark记录显示其再次被询问，从而验证了LRU逻辑。

## Task 5: Least Traffic Volume

### Coding

```
1 # add an empty heapq to store forwarding rules
2 # each element of tab should be a tuple like: (traffic, host, intf) --
  >traffic means the number of packets related
3 # when a heapq consists of tuples, it is organized based on the tuples' first
  elements
4 tab = []
```

```
1 # recording section
2 # src already recorded
3 if packet[0].src in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].src:
8             cur = rule
9             break
10    # adjust the heapq
11    # first delete, then add a new one since it's a tuple
12    new_rule = (cur[0], cur[1], input_port)
13    tab.remove(cur)
14    tab.append(new_rule)
15    heapq.heapify(tab)
```

```

16 # src not recorded yet
17 else:
18     # forwarding table is full
19     if len(tab) == max_rules:
20         heapq.heappop(tab)
21     tab.append((0, packet[0].src, input_port))
22     heapq.heapify(tab)

```

```

1 # forwarding section
2 # dst already recorded
3 if packet[0].dst in [rule[1] for rule in tab]:
4     # find the corresponding iterator for the current host
5     cur = iter(tab)
6     for rule in tab:
7         if rule[1] == packet[0].dst:
8             cur = rule
9             break
10    log_debug ("Flooding packet {} to {}".format(packet, cur[2]))
11    net.send_packet(cur[2], packet)
12    # add traffic volume and adjust the heapq
13    # first delete, then add a new one since it's a tuple
14    new_rule = (cur[0]+1, cur[1], cur[2])
15    tab.remove(cur)
16    tab.append(new_rule)
17    heapq.heapify(tab)
18 # dst not recorded yet
19 else:
20     # do nothing if dst is the switch itself
21     if packet[0].dst not in mymacs:
22         for intf in my_interfaces:
23             if input_port != intf.name:
24                 log_debug ("Flooding packet {} to {}".format(packet,
25 intf.name))
26                 net.send_packet(intf.name, packet)

```

## Testing

(此处提供的Test不全面)



Results for test scenario switch tests: 8 passed, 0 failed, 0 pending

Passed:

- 1 An Ethernet frame with a broadcast destination address should arrive on eth1
- 2 The Ethernet frame with a broadcast destination address should be forwarded out ports eth0 and eth2
- 3 An Ethernet frame from 20:00:00:00:00:01 to 30:00:00:00:00:02 should arrive on eth0
- 4 Ethernet frame destined for 30:00:00:00:00:02 should arrive on eth1 after self-learning
- 5 An Ethernet frame from 20:00:00:00:00:03 to 30:00:00:00:00:03 should arrive on eth2
- 6 Ethernet frame destined for 30:00:00:00:00:03 should be flooded on eth0 and eth1
- 7 An Ethernet frame should arrive on eth2 with destination address the same as eth2's MAC address
- 8 The switch should not do anything in response to a frame arriving with a destination address referring to the switch itself.

## Deploying

server1:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
2	0.100596398	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01
3	0.523463791	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x120d, seq=1/256, ttl=64 (reply in 4)
4	0.625066689	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x120d, seq=1/256, ttl=64 (request in 3)
5	1.046257343	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x120d, seq=2/512, ttl=64 (reply in 6)
6	1.149830414	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x120d, seq=2/512, ttl=64 (request in 5)
7	5.964527079	Private_00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.1
8	6.380728832	30:00:00:00:00:01	Private_00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01
9	12.084815042	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.2? Tell 192.168.100.3
10	38.318007216	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x1211, seq=1/256, ttl=64 (reply in 11)
11	38.418618476	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x1211, seq=1/256, ttl=64 (request in 10)
12	39.282417178	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x1211, seq=2/512, ttl=64 (reply in 13)
13	39.382510186	192.168.100.1	192.168.100.3	ICMP	98	Echo (ping) reply id=0x1211, seq=2/512, ttl=64 (request in 12)
14	43.383744011	30:00:00:00:00:01	Private_00:00:01	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
15	43.484072978	Private_00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.1 is at 10:00:00:00:00:01

server2:

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.1? Tell 192.168.100.3
2	12.084815594	30:00:00:00:00:01	Broadcast	ARP	42	Who has 192.168.100.2? Tell 192.168.100.3
3	12.185314046	20:00:00:00:00:01	30:00:00:00:00:01	ARP	42	192.168.100.2 is at 20:00:00:00:00:01
4	12.614372160	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x120e, seq=1/256, ttl=64 (reply in 5)
5	12.729434589	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x120e, seq=1/256, ttl=64 (request in 4)
6	13.148472439	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x120e, seq=2/512, ttl=64 (reply in 7)
7	13.249902982	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x120e, seq=2/512, ttl=64 (request in 6)
8	17.989440732	20:00:00:00:00:01	30:00:00:00:00:01	ARP	42	Who has 192.168.100.3? Tell 192.168.100.2
9	18.429659221	30:00:00:00:00:01	20:00:00:00:00:01	ARP	42	192.168.100.3 is at 30:00:00:00:00:01
10	29.184384162	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x1210, seq=1/256, ttl=64 (reply in 11)
11	29.285198949	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x1210, seq=1/256, ttl=64 (request in 10)
12	38.061852196	192.168.100.3	192.168.100.2	ICMP	98	Echo (ping) request id=0x1210, seq=2/512, ttl=64 (reply in 13)
13	38.161972749	192.168.100.2	192.168.100.3	ICMP	98	Echo (ping) reply id=0x1210, seq=2/512, ttl=64 (request in 12)
14	38.318007988	192.168.100.3	192.168.100.1	ICMP	98	Echo (ping) request id=0x1211, seq=1/256, ttl=64 (no response found!)

在client的cli中输入以下命令：

```
1 ping -c 2 192.168.100.1
2 ping -c 2 192.168.100.2
3 ping -c 2 192.168.100.2
4 ping -c 2 192.168.100.1
```

首先，设置最多容纳的rule数目为2。

在最开始的client对server1的两次请求中，结果与之前task一致。

之后，进行若干次client对server2的请求，增加server2的traffic\_volume，以达到清除有关server1对应端口信息的目的。

最后，再次进行client对server1的请求，server2的wireshark记录显示其再次被询问，从而验证了Traffic逻辑。

## 总结与感想

---

这次实验总体的代码量其实不算多，整体逻辑也比较清晰，但可能是因为对Python不熟悉的原因，很多时候想要实现一个feature却不知道该用什么合适的内置数据结构和方法，于是还要上网查阅，相对比较费时间。不过，从另一个角度来看，这也许是计网实验给我的另一个机遇，毕竟掌握一门语言最好的方式就是实战运用。

如果说lab\_1还是熟悉实验环境，这次就算是小试牛刀了吧。对相关工具的运用也有了进步，本来只是跟着教程重复，现在可以比较深入地理解其内部逻辑，甚至自己主动去探索进一步的用法。

最后，祝愿我们的Q&A小园地能越办越好，方便大家讨论的同时也可能造福以后的学弟学妹！！

## 总结与感想

---